

AMD Confidential Computing Technologies Evaluation Report

Roberto Castellotti
TU Munich

Abstract

We report a preliminary performance evaluation of AMD SEV (Secure Environment Virtualization) technologies. We focus on virtual machines, we are using QEMU/KVM as Hypervisor and we are running them on a powerful machine we control. Details about our environment can be found in section 4

Confidential Computing technologies may be predominant in the future, as more and more customers with sensitive computing workloads move their code from on-premise hardware to public cloud vendors. Being able to identify the performance bottlenecks may be crucial. After thoroughly explaining how the technologies introduced by AMD work we run some benchmarks to measure the impact these have on micro-benchmarks and traditional workloads such as compilation of popular open-source projects.

1 Introduction

Confidential Computing is a topic that started to become relevant in the last years. In the last two decades the way software is shipped to production changed radically, the majority of code deployed is hosted by cloud providers (Google Cloud Platform, Amazon Web Services, Microsoft Azure, etc.). Some code is still run on-premise for privacy/safety reasons. Confidential Computing aims to provide a safe environment for developers to run highly sensitive code on machines they don't own. It is logic customers want to be sure no one can access their disks, memory or CPU registers, neither other customers running virtual machines on the same hardware, nor whoever is controlling the hypervisor, be it the cloud vendor or, in worst case scenarios, malign actors who compromised the physical machines. Encryption at rest, designed to prevent the attacker from accessing the unencrypted data by ensuring the data is encrypted when on disk, has been around for a long time, and is currently supported by all major providers [15], [14], [2] but leaves a big part of daily computing unencrypted, namely RAM

and CPU registers, to tackle this issue major chip producers started to develop technologies to enable "confidential computing", namely AMD Secure Encrypted Virtualization (SEV) [5], Intel Trusted Domain Extensions (TDX) [10] and Arm Confidential Compute Architecture (CCA) [13]

2 Background

2.1 AMD Secure Memory Encryption (SME)

AMD Secure Memory Encryption (SME) is the basic building block for the more sophisticated technologies we'll cover later, so it is imperative we understand how it works. In a machine with SME enabled memory operations are performed via dedicated hardware, an entirely different chip on die. AMD EPYC™ introduced two hardware security components:

AES-128 hardware encryption engine embedded in memory controller, makes sure data to main memory is encrypted during write operations and decrypted during read operations, this memory controller is inside the EPYC SOC, so memory lines leaving the soc are encrypted

AMD Secure Processor a small processor providing cryptographic functionality for secure key generation and key management

The key used to encrypt and decrypt memory is generated securely by the AMD Secure-Processor (SMD-SP), a 32 bit micro-controller and it is not accessible by software running on the main CPU, furthermore SME does not require software running on main CPU to participate in Key Management making the enclave more secure.

The C-bit is a bit present in any memory page and indicates whether the current page is to be encrypted, it can be retrieved together with some additional information by running the `cpuid` command to inspect leaf `0x8000001F`, as specified by AMD Reference Manual (do i need to cite it?):

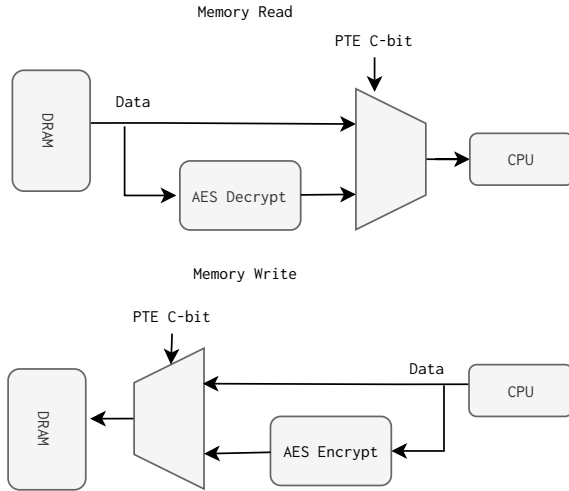


Figure 1: Memory Encryption Behavior, from [5]

```
$ cpuid -1 -1 0x8000001f
CPU:
AMD Secure Encryption (0x8000001f):
SME: secure memory encryption support = true
SEV: secure encrypted virtualize support = true
VM page flush MSR support = true
SEV-ES: SEV encrypted state support = true
SEV-SNP: SEV secure nested paging = true
VMPL: VM permission levels = true
Secure TSC supported = true
virtual TSC_AUX supported = false
hardware cache coher across enc domains = true
SEV guest exec only from 64-bit host = true
restricted injection = true
alternate injection = true
full debug state swap for SEV-ES guests = true
disallowing IBS use by host = true
VTE: SEV virtual transparent encryption = true
VMSA register protection = true
encryption bit position in PTE = 0x33 (51)
physical address space width reduction = 0x5 (5)
number of VM permission levels = 0x4 (4)
number of SEV-enabled guests supported = 0x1fd (509)
minimum SEV guest ASID = 0x80 (128)
```

SME is a very powerful mechanism to provide memory encryption, but it requires support from the Operating System/Hypervisor, **Transparent SME (TSME)** is a solution to encrypt every memory page regardless of the C-bit, as the name suggests this technology provides encryption without further modification to OS/HV, this may be crucial because Operating System developers don't have to support it and older operating systems can be run with TSME.

We now switch our focus to AMD SEV, a technology powered by AMD SME that enables Confidential Computing for virtual machines.




Figure 2: SEV security layers, from [5]

2.2 AMD Secure Encrypted Virtualization (SEV)

AMD SEV is an attempt to make virtual machines more secure to use by encrypting data to and from a virtual machine, and enables a new security model protecting code from higher privileged resources, such as hypervisors or some privileged code running on the physical machine hosting the virtual machines. In this context, as mentioned before, we should never trust the hypervisor since it may be compromised or acting maliciously by default.

SEV is an extension to the AMD-V architecture, when SEV is enabled SEV machines tag data with VM ASID (an unique identifier for that specific machine), this tag is used inside the SOC and prevents external entities to access it, when data leaves the chip we have no such problem because it is encrypted using the previously exchanged AES-128 bit key. The aforementioned expedients provide strong cryptography isolation between VMs run by the same hypervisor and between VMs and the hypervisor by itself. SEV guests can choose which pages to encrypt, this is handled setting the c-bit as mentioned before for SME. Only pages meant for outside communications are considered shared and thus not encrypted.

3 AMD Secure Encrypted Virtualization-Encrypted State (SEV-ES)

Up until now we only discussed encryption for memory, but a crucial portion of the system we want to protect are CPU registers, AMD SEV-ES encrypts all CPU register contents when a VM stops running. What this means is a malevolent actor is not able to read CPU's register contents when the machine is shutdown no matter the privilege level he acquired before, CPU register's state is saved and encrypted when the machine is shutdown.

Protecting CPU register may be a daunting task because sometimes an Hypervisor may need to access VM CPU's register to provide services such as device emulation. These accesses must be protected, ES technology allows the guest VM to decide which registers are encrypted, in the same vein

a machine can choose which memory pages are to be encrypted via the C-bit.

SEV-ES introduces a single atomic hardware instruction: VMRUN, when this instruction is executed for a guest the CPU loads all registers, when the VM stops running (VMEXIT), register's state is saved automatically to back to memory. These instructions are atomic because we need to be sure no one can sneak into this process and alter it, ES guarantees in this way it is impossible to leak memory.

Whenever hardware saves register it encrypts them with the very same AES-128 key we mentioned before, furthermore the CPU computes an integrity-check value and saves it into memory not accessible by the CPU, on next VMRUN instruction this will be checked to ensure nobody tried to tamper register's state. For further information about external communication consult the white-paper [11] and AMD reference manual chapter 15 [8].

Similarly to AMD-SEV AMD-ES is completely transparent to application code, only the guest VM and the Hypervisor need to implement these specific features.

3.1 AMD Secure Encrypted Virtualization-Secure Nested Paging (SEV-SNP)

After the introduction of AMD-SEV an AMD-ES AMD decided to introduce the next generation of SEV called Secure Nested Paging (SEV-SNP), this technology build on top of the aforementioned technologies and extends them further to implement strong memory integrity protection to prevent Hypervisor based attacks, such as **replay attacks** and **memory remapping, data corruption** and **memory aliasing**

replay attacks a malicious actor captures a state at a certain moment and modifies memory successfully with those values

data corruption even though an attacker cannot read a memory he can simply corrupt the memory to trick the machine into unpredicted behavior

memory aliasing an external actor may map a memory page to multiple physical pages

memory remapping the intruder maps a page to a different physical page

These attacks are a problem because a running program has no notion of memory integrity, they could end up in a state that was not originally considered by the developers and this may lead to huge security issues.

The basic principle of **SEV-SNP** integrity is that if a VM is able to read a private (encrypted) page of memory, it must always read the value it last wrote. (cite) What this means is the VM should be able to throw an exception if the memory a process is trying to access was tampered by external actors.

Threat Model

In this computing model we consider:

- **AMD System-On-Chip (SOC) hardware, AMD Secure Processor (AMD-SP)** and the **VM** are fully trusted, to this extend the VM should enable Full Disk Encryption (FDE) at rest, such as LUKS (cite), major cloud providers have been supporting FDE for long time: <https://cloud.google.com/docs/security/encryption/default-encryption>
- **BIOS** on the host system, the **Hypervisor, device drivers** and **other VMS** are fully untrusted, this means the threat model assumes they are malicious and may conspire to compromise the security of our Confidential Virtual Machine.

The way SEV-SNP ensures protection against the attacks we mentioned before is by introducing a new data structure, a **Reverse Map Table (RMP)** that tracks owners of memory pages, in this way we can enforce that only the owner of a certain memory page can alter it. A page can be owned by the VM, the Hypervisor or by the AMD Secure Processor. The RMP is used in conjunction with standard x86 page tables mechanisms to enforce memory restrictions and page access rights. RMP fixes replay, remapping and data corruption attacks.

RMP checks are introduced for write operations on memory, however external (Hypervisor) read accesses do not require them because we have AES encryption protecting our memory.

To prevent memory remapping a technique called **Page Validation** is introduced. Inside each RMP entry there is a Validated bit, pages assigned to guests that have no validated bit set are not usable by the Hypervisor, the guest can only use the page after setting the validated bit through a PVALIDATE instruction. The VM will make sure that it is not possible to validate a SPA (system physical address) corresponding to a GPA (Guest Physical Address) more than once.

More details are discussed in the introductory white-paper [6]

We now introduced every part of AMD's effort to popularize Confidential Computing, now we will proceed by giving instructions to start such machines using QEMU/KVM and we will run some benchmarks to measure how these technologies impact performance.

4 Environment

We are running our experiments in QEMU/KVM virtual machines, we are assigning 16GB of RAM and 16 vCPUs each, check the table 4 to see more details about our hardware and software versions.

QEMU is a generic open source machine emulator and virtualizer, we will use QEMU together with KVM, the Kernel Virtual machine to virtualize our machines.

OVMF is a project maintained by TianoCore aiming to enable UEFI support for virtual machines, it is based on EDK 2, we will use OVMF to generate the executable firmware and the non-volatile variable store, it is important to create a vm-specific copy of OVMF_vars.fd because the variable store should be private for every virtual machine. UEFI support is mandatory to run a SEV-SNP machine.

5 Benchmarks

We expect the aforementioned encryption techniques used in SEV/SEV-ES/SEV-SNP cause a degradation in performance, especially for memory-intensive workloads, additionally we want to investigate whether this technology slows down CPU intensive workloads. To quantify the price we have to pay to have a more secure virtual environment we are running three distinct kinds of benchmarks: Compilation, Memory Benchmarks and I/O benchmarks. We are using (<https://github.com/rcastellotti/tinyben>) [3], an experimental benchmarking tool aimed to replace the popular Phoronix Test Suite [12] benchmarking suite.

Compilation Benchmarks We are using compilation as an all-around benchmark because it is a "real world" benchmark. We are compiling some popular open-source projects like **Godot Game Engine**, **Linux** (defconfig), the entire **LLVM Project** (using ninja) and measuring how long does it take to complete the compilation process.

Memory Benchmarks To benchmark memory we are using ssvb/tinymembench [16], a tool to measure memory throughput and latency, we are mainly interested in bandwidth for the MEMCPY and MEMSET operations

I/O Benchmarks To measure I/O performance we are running two different benchmarks, first we are performing 2500 sqlite insertions in a table, then we are running redis-benchmark, a tool shipped together with redis. We are interested in Requests per Second and latency for the main operations (namely SET and GET)

As we can see in table 5 the overhead introduced by the usage of Confidential Computing technology across the three categories of benchmarks is generally mild, and negligible in almost every benchmark, with the exception for LZ4-compression. further investigation tweaking the amount of resources available to each machine is necessary to correctly esteem the exact performance degradation values.




Figure 3: Compilation benchmarks, in milliseconds




Figure 4: tinymembench benchmarks, MEMSET and MEMCPY

One reason why CPU-intensive workloads are little to no penalized by SEV-SNP is the processor is not doing much more when executing that kind of workload on a confidential machine, the instructions that really slow down are VMEXIT and VMRUN, we are not issuing any of those during the compilation of a certain program. This is good news, it means we don't have to sacrifice performance to have a safer environment.

We ran some additional benchmarks regarding the different storage virtualization techniques, more specifically we tested **virtio-scsi**, **virtio-blk** and **nvme** technologies. Virtio is the main platform for IO virtualization in KVM, essentially it provides a common framework for hypervisors to do IO virtualization.

nvme is the user-space NVMe driver that enables virtual machines to interact with NVMe devices

virtio-blk devices are very simple block devices, the front-end driver reads and writes by appending commands to the virtualization queue so that the back-end driver can process them on the host.

Host CPU	AMD EPYC 7713P 64-Cores
Host Memory	HMAA8GR7AJR4N-XN (Hynix) 3200MHz 64 GB × 8 (512GB)
Host Kernel	6.3.0-rc2 #1-NixOS SMP PREEMPT_DYNAMIC (NixOS 23.05) commit: fea9b78
QEMU	8.0.0 (AMD) (patched) commit: a248931
OVMF	Stable 202211 (patched) commit: 6598f62
Guest vCPUs	16
Guest Memory	16GB
Guest Kernel	5.19.0-41-generic #42-Ubuntu SMP PREEMPT_DYNAMIC (Ubuntu 22.10)

Table 1: Experiment environment

benchmark	SEV-SNP	NOSEV
Godot Game Engine compilation (LIB)	410 sec	395 sec
ImageMagick compilation (gcc) (LIB)	70 sec	69 sec
Linux compilation (defconfig) (LIB)	329 sec	322 sec
LLVM Project compilation (ninja) (LIB)	774 sec	733 sec
Sqlite 2500 insertions (LIB)	4.106 sec	4.31 sec
ssvb/tinymembench MEMCPY (HIB)	22727.0 MB/s	23628.2 MB/s
ssvb/tinymembench MEMSET (HIB)	33170.8 MB/s	36367.7 MB/s
redis-benchmark SET RPS (HIB)	91911.76	110864.74
redis-benchmark SET average latency (LIB)	0.283 msec	0.234 msec
redis-benchmark GET RPS (HIB)	94339.62	110987.79
redis-benchmark GET average latency (LIB)	0.275 msec	0.233 msec

Table 2: Comparing benchmark results run on two QEMU vms (16 vCPUs, 16GB ram, virtio-scsi storage). HIB and LIB mean respectively "Higher is Better" and "Lower is Better", RPS stands for "Requests Per Second"




Figure 5: Miscellaneous benchmarks, lz4 compression and decompression and Sqlite insertions

virtio-scsi aims to overcome some limitations introduced by virtio-blk, support more devices per guest (one PCI device per disk is not a limiting factor anymore) and supports technologies like multiqueueing while keeping the performances of virtio-blk, additionally virtio-scsi provides a pass-through technology to present physical storage devices directly to guests

We perform some measurements using the FIO [1] benchmarking tool, we use the same measurements used in [17] to evaluate: bandwidth, average latency and IOPS, we use the `-direct=1` flag to use unbuffered I/O.

The benchmarks highlight some interesting patterns: average latency is barely impacted, while bandwidth is severely impacted, especially in read workloads, this might be related to the usage of the aforementioned **Reverse Map Table**.

6 Future Work

Further performance evaluation for SEV/SEV-ES/SEV-SNP confidential virtual machines will include:

- repeating micro-benchmarks and FIO benchmarks tweaking machines configuration (vCPUs, RAM, storage virtualization)
- analyzing IO performance when using SEV Trusted Input Output [7]

benchmark	test cases	FIO Configuration (bs, rw, iodepth, numjobs)
bandwidth	read	(128K, read, 128, 1)
bandwidth	write	(128K, write, 128, 1)
IOPS	randread	(4K, randread, 32, 4)
IOPS	mixread	(4K, randread 70%, 32, 4)
IOPS	mixwrite	(4K, randwrite 30%, 32, 4)
IOPS	randwrite	(4K, randwrite, 32, 4)
average latency	randwrite	(4K, randread, 1, 1)
average latency	randwrite	(4K, randwrite, 1, 1)
average latency	read	(4K, read, 1, 1)
average latency	write	(4K, write, 1, 1)

Table 3: FIO benchmarks, from [17]




Figure 6: Storage Bandwidth Benchmarks




Figure 8: IOPS benchmarks




Figure 7: Average Latency Benchmarks

mode	storage	nosev	sev	result
read	blk	2068866712	829144265	0.400772
read	nvme	2255760134	1483068817	0.657459
read	scsi	2644684295	732429620	0.276944
write	blk	2010752479	1398101333	0.695312
write	nvme	1664716006	1518729595	0.912306
write	scsi	1913978295	1416545941	0.740106

Table 4: FIO bandwidth benchmarks results, measured in KB/s, ratio is the ratio between **sev** and **nosev**

mode	storage	nosev	sev	result
mixread	blk	24324.394544	19608.347670	0.806119
mixread	nvme	16936.555110	16221.782178	0.957797
mixread	scsi	23418.259782	21283.104652	0.908825
mixwrite	blk	22876.690811	19261.131521	0.841954
mixwrite	nvme	15671.907694	13144.662288	0.838740
mixwrite	scsi	22213.710702	18586.500284	0.836713
randread	blk	24786.686838	19354.991140	0.780862
randread	nvme	16517.169681	13646.936332	0.826227
randread	scsi	24944.714055	20894.627770	0.837637
randwrite	blk	23045.626374	18642.013938	0.808918
randwrite	nvme	16179.730897	14074.845638	0.869906
randwrite	scsi	22775.325804	19140.186916	0.840391

Table 5: FIO IOPS benchmarks results, ratio is the ratio between **sev** and **nosev**

mode	storage	nosev	sev	ratio
randread	blk	36732.676746	44428.919037	1.209520
randread	nvme	43056.001194	43182.237293	1.002932
randread	scsi	37267.866497	44158.765942	1.184902
randwrite	blk	39318.967918	42268.305611	1.075011
randwrite	nvme	40921.898750	48371.563309	1.182046
randwrite	scsi	40428.971912	47231.466854	1.168258
read	blk	34110.888638	40344.405910	1.182743
read	nvme	39693.322182	37262.537930	0.938761
read	scsi	35575.824402	40626.232517	1.141962
write	blk	37879.406013	41905.050934	1.106275
write	nvme	51967.809612	46775.977871	0.900095
write	scsi	40000.623734	45465.357731	1.136616

Table 6: FIO average latency benchmarks, measured in nanoseconds, ratio is the ratio between **sev** and **nosev**

- analyzing whether having multiple SEV machines on the same physical host incurs in a large overhead
- analyze performance overhead caused by VMRUN and VMEXIT when SEV-ES is enabled.
- exploring Confidential Containers [4] project to identify what is worth benchmarking
- understand AMD Secure VM Service Module [9], a module to offload sensitive operations onto a privileged guest and see if this leads to a performance improvement

References

- [1] Jens Axboe. *Flexible I/O tester rev. 3.35*. URL: https://fio.readthedocs.io/en/latest/fio_doc.html (visited on 07/07/2023).
- [2] Microsoft Azure. *Azure Data Encryption at rest*. URL: <https://learn.microsoft.com/en-us/azure/security/fundamentals/encryption-atrest> (visited on 07/07/2023).
- [3] Roberto Castellotti. *a <1000 loc benchmark suite in python*. URL: <https://github.com/rcastellotti/tinyben> (visited on 07/07/2023).
- [4] Confidential Containers. *Confidential Containers*. URL: <https://github.com/confidential-containers> (visited on 07/21/2023).
- [5] Tom Woller David Kaplan Jeremy Powell. *AMD Memory Encryption*. 2021. URL: <https://www.amd.com/system/files/TechDocs/memory-encryption-white-paper.pdf> (visited on 07/07/2023).
- [6] Advanced Micro Devices Inc. *AMD SEV-SNP: Strengthening VM Isolation with Integrity Protection and More*. 2020. URL: <https://www.amd.com/system/files/TechDocs/SEV-SNP-strengthening-vm-isolation-with-integrity-protection-and-more.pdf> (visited on 07/07/2023).
- [7] Advanced Micro Devices Inc. *AMD SEV-TIO: Trusted I/O for Secure Encrypted Virtualization*. 2023. URL: <https://www.amd.com/system/files/documents/sev-tio-whitepaper.pdf> (visited on 07/07/2023).
- [8] Advanced Micro Devices Inc. *AMD64 Architecture Programmer's Manual Volumes 1–5*. 2023. URL: <https://www.amd.com/system/files/TechDocs/40332.pdf> (visited on 07/07/2023).

- [9] Advanced Micro Devices Inc. *Secure VM Service Module for SEV-SNP Guests*. 2022. URL: <https://www.amd.com/system/files/TechDocs/58019-svsm-draft-specification.pdf> (visited on 07/21/2023).
- [10] Intel®. *Intel® Trust Domain Extensions*. URL: <https://cdrdv2.intel.com/v1/dl/getContent/690419> (visited on 07/07/2023).
- [11] David Kaplan. *Protecting VM Register State with SEV-ES*. 2017. URL: <https://www.amd.com/system/files/TechDocs/Protecting%20VM%20Register%20State%20with%20SEV-ES.pdf> (visited on 07/07/2023).
- [12] Michael Larabel. *The Phoronix Test Suite open-source, cross-platform automated testing/benchmarking software*. URL: <https://github.com/phoronix-test-suite/phoronix-test-suite> (visited on 07/07/2023).
- [13] Arm Limited. *Introducing Arm Confidential Compute Architecture*. URL: <https://developer.arm.com/documentation/den0125/latest> (visited on 07/07/2023).
- [14] Google Cloud Platform. *Default encryption at rest*. URL: <https://cloud.google.com/docs/security/encryption/default-encryption> (visited on 07/07/2023).
- [15] Amazon Web Services. *Encryption of Data at Rest*. URL: <https://docs.aws.amazon.com/whitepapers/latest/efs-encrypted-file-systems/encryption-of-data-at-rest.html> (visited on 07/07/2023).
- [16] Siarhei Siamashka. *Simple benchmark for memory throughput and latency*. URL: <https://github.com/ssvb/tinymembench> (visited on 07/07/2023).
- [17] Shuai Xue et al. “Spool: Reliable Virtualized NVMe Storage Pool in Public Cloud Infrastructure”. In: *Proceedings of the 2020 USENIX Conference on Usenix Annual Technical Conference*. USENIX ATC’20. USA: USENIX Association, 2020. ISBN: 978-1-939133-14-4.

Appendix

A A simple attack to demo Confidential Computing

Let’s demo a very simple attack, first of all we start two machines, *sev* and *nosev*, the former has SEV-SNP enabled, as we can check:

```
sev:~$ sudo dmesg | grep SEV
Memory Encryption Features active: AMD SEV SEV-ES SEV-SNP
SEV: Using SNP CPUID table, 31 entries present.
SEV: SNP guest platform device initialized.
```

We will write something into a file and cat it in order to load the data in memory

```
sev:~$ echo "hi from SEV!" > sev.txt
sev:~$ cat sev.txt
hi from SEV!
```

```
nosev:~$ echo "hi from NOSEV!" > nosev.txt
nosev:~$ cat nosev.txt
hi from NOSEV!
```

Now we can dump the memory for the processes using *gcore*

```
h:~$ sudo gcore -o mem-dump <SEV_PID>
h:~$ grep -rnw mem-dump.<SEV_PID> -e "hi from SEV!"
h:~$
```

```
h:~$ sudo gcore -o mem-dump <NOSEV_PID>
h:~$ grep -rnw mem-dump.<NOSEV_PID> -e "hi from NOSEV!"
grep: mem-dump.<NOSEV_PID>: binary file matches
```

From the host machine we are able to see *nosev*’s machine memory while this is not possible with SEV enabled.